

11-755: MLSP

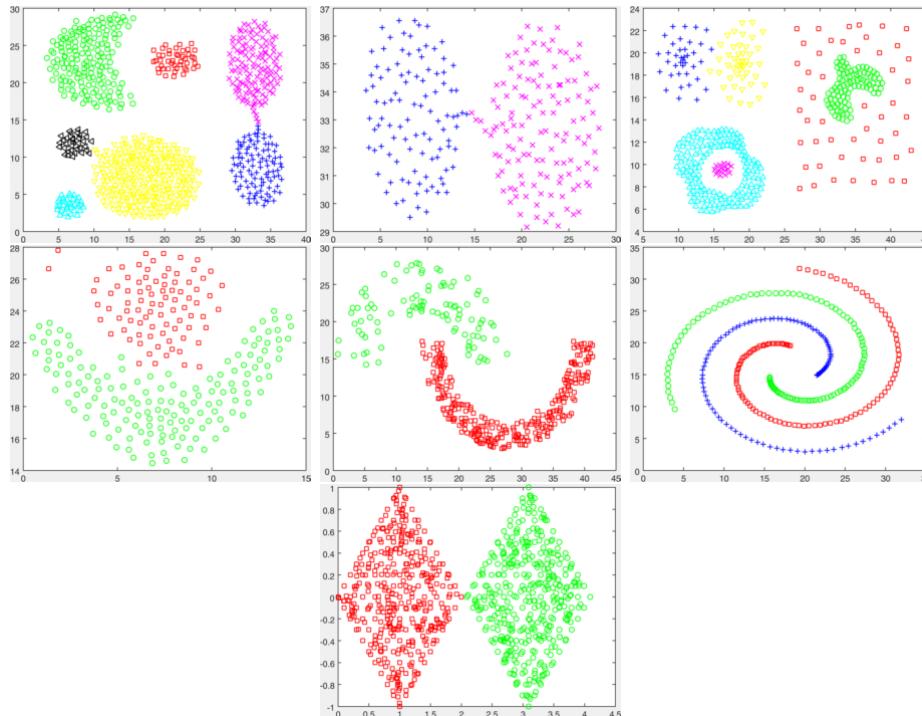
HW #3: Clustering and EM

Nikolas Wolfe

December 15, 2015

Problem 1: K-Means & Spectral Clustering

You are given a number of toy datasets. The visualized ground truth clustering configuration of these data: Aggregation (K=7), Bridge (K=2), Compound (K=6), Flame (K=2), Jain (K=2), Spiral (K=3) and TwoDiamond (K=2) are shown as follows:



Answer:

I implemented this code in Java, using the Efficient Java Matrix Library (EJML) package to do matrix multiplication and Eigen/Singular Value decomposition. In general, spectral clustering works a lot better than K-means and is able to determine a lot of structure (such as the spiral pattern) that k-means is unable to find.

In terms of evaluation, I ran the clustering algorithms and generated the confusion matrix between the most common label that emerges in each cluster and their actual labels. A diagonal matrix in this case is of course perfect performance. The accuracy can be determined from looking at the counts of labeled points for each class.

Dataset	Optimal K	Group Counts
Aggregation	7	$[K_1, K_2, K_3, K_4, K_5, K_6, K_7] = [45, 170, 102, 273, 34, 130, 34]$
Compound	6	$[K_1, K_2, K_3, K_4, K_6] = [50, 92, 38, 45, 158, 16]$
Spiral	3	$[K_1, K_2, K_3] = [101, 105, 106]$
Bridge	2	$[K_1, K_2] = [102, 13]$
Flame	2	$[K_1, K_2] = [87, 153]$
Jain	2	$[K_1, K_2] = [276, 97]$
Two Diamonds	2	$[K_1, K_2] = [400, 400]$

With this information, we can directly evaluate the quality of the Spectral and K-means clustering algorithms by simply comparing the diagonal of their confusion matrices (most common group label vs. actual label) to the vectors listed above. To run the java package, simply run the following from the root of this project:

```
java -jar nwolfe_hw3.jar
```

This will run the main method in the class `ClusterDriver.java`, which will then iterate through a small range of σ values which are known to produce generally good spectral clusters, though it is often the case that several random restarts are required in order for the spectral clustering algorithm to settle into the optimal cluster groupings. In this report the best clusters have been chosen, though it is the case that when *you* run it, it might produce an optimal grouping.

This code also produces the results as expected by the provided MATLAB scripts as CSV files in the root directory in which the jar is run. To generate the visualizations of the data, you may select the individual files

you want and use the `visualize.m` script as provided. Alternatively, you may use the results I have included in the report and simply go to the `external` folder and run the `gen_all_figs.m` script.

```
cd external/  
matlab gen_all_figs.m
```

Alternatively, you could just look at the pictures in this here report. That's probably easier.

Questions

In your report, answer the following questions:

1. Is the objective function of k-means a convex one?

Answer: The objective function which we are trying to minimize is the sum of Euclidean distances from each cluster centroid to each cluster datapoint. However, the number of clusters K is discrete, which means we cannot take the derivative of this function and determine the optimal grouping (i.e. the minimum distance sum over all clusters) either analytically or through gradient descent or some other minimization technique. (This is also the case for spectral clustering.) Furthermore, the 2nd derivative of the Euclidean distance function is not zero. Because there is no way to find the optimal grouping (and perhaps no grouping actually exists), this is not a convex optimization problem.

2. Can the iterative updating scheme in k-means return a solution which achieves the global minimum of the k-means objective function?

Answer: Newp. No. Sorry! As stated above, there is no guaranteed minimum for the objective function because it isn't convex. There may be an infinite number of viable ways to cluster the data.

3. If yes, why? If not, why not and are there any practical ways to partially fix this?

Answer: One of the ways to fix this is, first off, to attempt to pick the initialization of the cluster centroids according to some logic. In this case we happen to know the optimal value of K from the labeled data. This may result in a better teasing apart of the datapoints and less moving around before the algorithm converges. Another way might be to use bottom up or top down clustering, i.e. to start with

a single cluster containing all the data, perturb the centroid and then continually break the clusters apart until an optimal value of K is reached. Or we could start with a large number of clusters and then continually merge them according to some criteria until we reach a desired value of K .


Results

For each problem I have printed the output of the algorithm after a single run. At the beginning of each section is a comparison of 3 pictures.

Aggregation

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Aggregation: Spectral Clustering

```
Reading data values in data/Aggregation.csv...
Reading data values in data/Bridge.csv...
Reading data values in data/Compound.csv...
Reading data values in data/Flame.csv...
Reading data values in data/Jain.csv...
Reading data values in data/Spiral.csv...
Reading data values in data/TwoDiamonds.csv...
```

```
===== AGGREGATION: SPECTRAL CLUSTERING =====
With sigma = 0.47
```

label	1	2	3	4	5	6	7
1	45	0	0	0	0	0	0
2	0	170	0	0	0	0	0
3	0	0	1	0	0	0	0
4	0	0	0	270	0	0	0
5	0	0	0	0	34	0	0
6	0	0	101	0	0	130	0
7	0	0	0	3	0	0	34

With sigma = 0.48

label	1	2	3	4	5	6	7
1	45	0	0	0	0	0	0
2	0	170	0	0	0	0	0
3	0	0	102	0	0	2	0
4	0	0	0	270	0	0	0
5	0	0	0	0	34	0	0
6	0	0	0	0	0	128	0
7	0	0	0	3	0	0	34

With sigma = 0.49

... Output truncated ...

In the second run above, the diagonal is almost perfect, save for a few items misplaced in groups 3 and 7.

Aggregation: K-Means Clustering

===== AGGREGATION: K MEANS =====

label	1	2	3	4	5	6	7
1	0	0	0	0	0	0	0
2	17	170	0	0	0	0	0
3	0	0	102	0	0	3	0
4	0	0	0	273	34	0	34
5	0	0	0	0	0	0	0
6	28	0	0	0	0	127	0


```
7 | 0 0 0 0 0 0 0
```

In the k-means clustering it's clear that several groups did not even emerge from the data (though this does not mean they were empty.)

Bridge

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Bridge: Spectral Clustering

```
===== BRIDGE: SPECTRAL CLUSTERING =====
```

```
With sigma = 0.42
```

```
... Output truncated ...
```

```
With sigma = 0.44
```

```
label    1    2
-----
1 |      102  2
2 |       0 128
```

```
With sigma = 0.45
```

label	1	2
1	58	58
2	44	72

In one case the grouping was nearly perfect, and in another the data was more or less spread around a different kind of grouping.

Bridge: K-Means Clustering

===== BRIDGE: K MEANS =====


label	1	2
1	0	0
2	102	130

Again, K-means was not able to find the right clustering...

Compound

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Compound: Spectral Clustering

===== COMPOUND: SPECTRAL CLUSTERING =====

With sigma = 0.154

label	1	2	3	4	5	6
1	21	0	1	1	0	0
2	6	92	0	0	0	0
3	12	0	33	0	0	0
4	3	0	1	44	0	0
5	5	0	2	0	158	0
6	3	0	1	0	0	16

With sigma = 0.155

label	1	2	3	4	5	6
1	20	0	2	1	0	0
2	6	92	1	44	0	0
3	11	0	35	0	0	0
4	0	0	0	0	0	0
5	7	0	0	0	158	0
6	6	0	0	0	0	16

... Output truncated ...

This grouping is particularly interesting because the groups in the compound picture were very difficult to tease apart, and yet we see here a diagonal clearly taking shape, except for some overlap with group 1.

Compound: K-Means Clustering

===== COMPOUND: K MEANS =====


label	1	2	3	4	5	6
1	0	0	0	0	0	0
2	48	92	0	0	0	0
3	0	0	38	1	0	0

4		2	0	0	44	0	0
5		0	0	0	0	158	16
6		0	0	0	0	0	0

Flame

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Aggregation: Flame Clustering

===== FLAME: SPECTRAL CLUSTERING =====
 With sigma = 0.735

label	1	2	

1		0	0
2		87	153

Flame: K-Means Clustering

===== FLAME: K MEANS =====


label	1	2	

1		0	0

Jain

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Jain: Spectral Clustering

===== JAIN: SPECTRAL CLUSTERING =====

With sigma = 0.303

label	1	2
1	276	0
2	0	97

With sigma = 0.304

label	1	2
1	276	97
2	0	0

... Output truncated ...

Jain: K-Means Clustering


===== JAIN: K MEANS =====

label	1	2
1	276	97
2	0	0

Spiral

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



Spiral: Spectral Clustering

===== SPIRAL: SPECTRAL CLUSTERING =====

With $\sigma = 0.2$

label	1	2	3
1	101	0	0
2	0	105	0
3	0	0	106

With `sigma = 0.30000000000000004`

```
label  1      2      3
-----
1 |      101  0      0
2 |      0    105  0
3 |      0      0    106
```

With `sigma = 0.4`

```
label  1      2      3
-----
1 |      101  0      0
2 |      0    105  0
3 |      0      0    106
```

... Output truncated ...

Given the shape in this case we can really see the difference between spectral clustering and k-means. This grouping was found perfectly, all zeros in the off-diagonal elements, with even a variable number of settings for σ . Looking at the k-means result below, it's clear that there's a real distinction here.

Spiral: K-Means Clustering


===== SPIRAL: K MEANS =====

```
label  1      2      3
-----
1 |      66    54    53
2 |      0      0      0
3 |      35    51    53
```

Two Diamonds

From left to right, the images are:

1. Ground Truth
2. Spectral Clustering Result
3. K-Means Clustering Result



In the above case it appears that the two centroids in the k-means cluster were positioned in such a way that the points are ambiguously scattered, although I should note that I did not control for bad k-means results... I have only tried to present the best spectral clustering results because they're more fun anyway.

Two Diamonds: Spectral Clustering

===== TWO DIAMONDS: SPECTRAL CLUSTERING =====

With sigma = 20.0

label	1	2
1	400	0
2	0	400

With sigma = 21.0

label	1	2
1	204	198
2	196	202

... Output truncated ...

Two Diamonds: K-Means Clustering

===== TWO DIAMONDS: K MEANS =====

label	1	2
1	203	196
2	197	204

Questions & Conclusion

Compare your spectral clustering results with k-means. It is natural that on certain hard toy examples, both method won't generate perfect results? In your report, briefly analyze what is the advantage or disadvantage of spectral clustering over k-means. Why it is the case?

Answer:

It is natural in the sense that the Gaussian kernel function used in spectral clustering is good at separating certain kinds of data and not others. Clearly separable groups, even if they have complicated boundaries (such as the spiral) are

Overall, I found that SVD was better at finding the basis vectors for the projection of the points into k dimensional space. Furthermore in many cases it required a lot of random restarts before the groups coalesced into something reasonable.

Problem 2: Expectation Maximization

Let Z be the sum of two random variables X and Y , ($Z = X + Y$), where X and Y are drawn independently from below discrete probability distributions with probability mass functions defined as:

$$P(X = n) = P_1(1 - P_1)^{n-1} \quad (1)$$

$$P(Y = n) = P_2(1 - P_2)^{n-1} \quad (2)$$

Given samples of Z , derive an EM algorithm to estimate P_1 and P_2 .

Answer: We have two discrete probability distributions. Samples of Z are always the sum of X and Y , so Z is a multinomial distribution. We want to use the EM algorithm to estimate the maximum likelihood values of P_1 and P_2 that generated the samples of Z . We don't know P_1 and P_2

from observing Z , so we have to iteratively *fractionate* the observations and reestimate our expectation for the values of P_1 and P_2 .

There will be two steps in defining this EM algorithm: The E-step and the M-step. In the E-step, we compute an expectation for the data in Z for the current (or assumed) values of P_1 and P_2 . In the M-step, we treat this value as the ground truth and update the maximum likelihood values of P_1 and P_2 . Overall, we are attempting to optimize the auxiliary function:

$$Q(\theta, \theta') = \sum_H P(H|Z, \theta') \log(P(H, Z|\theta)) \quad (3)$$

H is the unseen variable (or variables)... In this case, H represents the two hidden variables X and Y , and Z is the observation. θ represents the current values of the estimates of the model parameters (i.e. P_1 and P_2), and θ' represents the values in the previous iteration of the algorithm.

This is just like the dice problem we solved in class. Given counts $c_1, c_2, c_3 \dots c_n$ for each discrete value of n that X and Y take on, we can first calculate the total log probability of P_1 and P_2 . (We will use log probability in this case because eventually we'll be taking the derivative to get the maximum likelihood value for P_1 and P_2 and the log is easily differentiable whereas this function as is totally sucks.)

The total log probabilities of X and Y are the sum of the probabilities of each value of n times the counts c_n . This is:

$$\log(P(X)) = \sum_n c_n \log(P_1) + \log(1 - P_1)(n - 1) \quad (4)$$

$$\log(P(Y)) = \sum_n c_n \log(P_2) + \log(1 - P_2)(n - 1) \quad (5)$$

To find the maximum likelihood values of P_1 and P_2 we differentiate this equation and set it to zero:

$$\frac{\partial}{\partial P_1} \log(P(X)) = \frac{\partial}{\partial P_1} \left(\sum_n c_n \log(P_1) + \log(1 - P_1)(n - 1) \right) \quad (6)$$

$$0 = \frac{\partial}{\partial P_1} \left(\sum_n c_n \log(P_1) + \log(1 - P_1)(n - 1) \right) \quad (7)$$

$$0 = \frac{c_n}{P_1} - \frac{n - 1}{(1 - P_1)} \quad (8)$$

$$\frac{c_n}{P_1} = \frac{n - 1}{(1 - P_1)} \quad (9)$$

$$P_1 = \frac{c_n}{c_n + n - 1} \quad (10)$$

Likewise for P_2 :

$$P_2 = \frac{c_n}{c_n + n - 1} \quad (11)$$

For the first estimates of P_1 and P_2 , we sum this value over all counts of n in Z . Unfortunately, Z represents the sum of two hidden variables X and Y , and we do not know what values they actually take. Because we do not know, we fractionate the value of Z into two bins, each of which can take on n values. As we already know:

$$Z = X + Y \quad (12)$$

Every Z *must* contain a discrete X and Y value. Thus the values of X and Y represent *fractions* of Z implicitly. By fractionating Z , we get possible X and Y values directly. X and Y are the hidden variables, dependent on P_1 and P_2 , and Z is the observation. From the perspective of just X , the log likelihood for a given n is:

$$\log(P(X = n|Z)) \quad (13)$$

Which can be computed using Bayes' Rule:

$$\log(P(X = n|Z)) = \log(P(Z|X = n)) + \log(P(X = n)) - \log(P(Z)) \quad (14)$$

$$\log(P(X = n|Z)) = \log(P(Z|X = n)) + \log(P(X = n)) + \text{Const} \quad (15)$$

We have $\log(P(X = n))$ from above:

$$\log(P(X = n)) = \log(P_1) + \log(1 - P_1)(n - 1) \quad (16)$$